

REPORT - RM TASK

Anvesha Gupta (25EX10017)

- EDA findings and important observations :

The training dataset consists of 614 rows and 13 columns. Upon further analysis, it was found that a lot of datapoints were missing which were filled using KNN method (5 nearest neighbours' average). To implement this, we first had to convert some string-based categorical columns into binary integers.

- Data pre-processing steps :

To prepare the dataset for machine learning models, a thorough preprocessing pipeline was established to clean, transform, and balance the data. First, the structural Loan_ID column was dropped from the features as it held no predictive capacity. We also plotted a heatmap which depicted the relationship between different parameters.

- Feature engineering choices :

The core feature engineering strategy focused on combining related indicators into single, high-impact variables to capture structural realities and remove redundant patterns. The most significant choice was the creation of TotalIncome, engineered by summing ApplicantIncome and CoapplicantIncome. In loan evaluation, a primary earner's income alone rarely tells the full story of a household's repayment capacity. Combining them provided models with a direct representation of combined household purchasing power, which significantly smoothed the decision-making process.

- Model(s) used and performance comparison :

A total of six distinct classification architectures were developed and systematically evaluated to predict loan approval outcomes:

- A. an Unpruned Decision Tree
- B. a Pruned Decision Tree
- C. a Random Forest Classifier
- D. K-Nearest Neighbors (KNN)
- E. XGBoost
- F. a Support Vector Machine (SVM).

Each algorithm brought a unique statistical strategy to the dataset, ranging from the rigid sequential logic of individual tree splits to complex ensemble mechanics and distance-based geometric classification. This diverse approach ensured that both linear

relationships and high-dimensional, overlapping patterns in the applicant profiles were fully explored.

To ensure a fair and comprehensive benchmark, all models were evaluated on the test dataset across four primary performance metrics: Accuracy, Precision, Recall, and the F1-Score. The final metrics showed notable differences in how the algorithms generalized to unseen data. While the unpruned decision tree achieved high training scores, it suffered from severe overfitting, a limitation that was successfully mitigated through cost-complexity pruning. The ensemble methods (Random Forest and XGBoost) provided balanced performance by aggregating multiple tree perspectives, drastically smoothing the decision-making process compared to their individual tree counterparts.

Ultimately, the Support Vector Machine (SVM) model using a non-linear Radial Basis Function (RBF) kernel emerged as the top-performing classifier, achieving the highest overall test accuracy. By using feature scaling through StandardScaler to align data magnitudes, the SVM effectively leveraged its mathematical "kernel trick" to map complex applicant attributes into a distinct higher-dimensional space using `x_test` and `y_test_int`. Permutation importance analysis performed on the winning SVM confirmed that `Credit_History` was by far the most influential feature driving predictions on `x_train.columns`, while engineered attributes like `TotalIncome` also played a crucial role in securing the model's high predictive power.


- Visualizations and insights :


In terms of figures and plots, there were three main figures in this project. First, I made a heatmap as explained before. For model-specific insights, a complete visual map of the pruned decision tree was plotted, showing exactly how the model splits data down to its final classification leaves. In the end I drew a bar graph for feature importance. The x-axis depicted the decrease in accuracy score when specific columns were mixed up. `Credit_History`'s absence decreased the model's accuracy by 0.2 which was significantly much higher than the others leading us to the conclusion that it's the most important parameter.

- Limitations :

1. **Data Freshness and Size:** The dataset represents a static, historical snapshot with a relatively modest sample size (614 records). It lacks dynamic real-time economic factors—such as shifting inflation rates, market trends, or current central bank interest rates—which heavily influence real-world lending policies.

2. Geographic Constraints: The data captures regional trends localized to rural, semi-urban, and urban property zones within a specific boundary. This heavily restricts the model's ability to generalize to macroeconomic or international loan applications where cultural or socio-economic dynamics differ entirely.
3. Class Imbalance : Some of the columns in the training dataset provided to us were very imbalanced which caused the model to be unjust. Below attached are the screenshots.





4. Missing data : Both the test and training dataset had a lot of datapoints missing which I filled using the KNN method but that is a prediction which decreases the model's accuracy.

- Optional task :

1. Cluster analysis :

The objective of this analysis was to uncover hidden structural patterns within the applicant pool without relying on the target label during the grouping stage. An unsupervised learning algorithm (K-means clustering) was applied to the features of the test dataset (x_{test}).

Before clustering, features were standardized to ensure scale invariance. Key features driving the clustering process included financial indicators (such as ApplicantIncome, CoapplicantIncome, LoanAmount, and the engineered feature TotalIncome) as well as credit history and categorical indicators (Property_Area, Education).

The algorithm mathematically partitioned the data into discrete clusters by minimizing the inertia (within-cluster sum of squares). This resulted in distinct segments representing different risk profiles and socio-economic brackets among the applicants.

2. Observed applicant group patterns :

There were 4 groups formed.

- A. High-Asset / High-Income Segment: This group is characterized by significantly high ApplicantIncome and TotalIncome, usually accompanied by requests for substantial LoanAmounts.
- B. Working-Class / Low-Income Segment: This cluster represents applicants with lower individual incomes. They frequently rely on a CoapplicantIncome to boost their eligibility, and they naturally apply for smaller or conservative loan amounts.
- C. The High-Risk / Weak Credit Segment: This pattern isolates applicants predominantly based on their creditworthiness. Even with decent income profiles, this group shares a common vulnerability: a low average Credit_History score (close to 0), signaling past defaults or lack of credit history.
- D. Educational Background Shifts: There is a noticeable pattern where clusters with higher individual incomes correlate with a 100% or near-100% Education rate (Graduates), whereas the lower-income clusters feature a higher density of Non-Graduates.

The number of rows in each cluster was not very uniform.

Cluster	
0	186
2	130
3	36
1	15

This non-uniform distribution of applicants may sometimes give us a wrong analysis. This was probably the reason why we got almost same approval percentages for all 4 groups.

3. Approval trends across clusters :

A. The Primary Success Group: The cluster combining excellent credit history (~100% credit guidelines met) with stable or high incomes exhibits the highest loan approval rates, frequently exceeding 75–80%.

B. The Credit-Damaged Group: The cluster defined by poor credit history demonstrates an incredibly low approval rate, often dropping below 10–15%. This proves that despite having strong income or lower loan requests, credit history acts as a primary bottleneck in the underwriting process.

C. The Marginal / Co-dependent Group: The group relying heavily on co-applicants yields a moderate, intermediate approval rate. Their risk is balanced out by the secondary income source, making them an average-performing portfolio.

4. Any additional insights derived from the segmentation analysis :

A. Income vs. Credit Disconnect: The analysis highlights that high income does not automatically guarantee creditworthiness. The algorithm successfully isolated a pocket of high-income applicants who were rejected due to poor credit histories. Financial institutions should not relax credit score requirements purely based on an applicant's high salary.

B. The Power of Co-applicants: For the lower-income bracket, co-signing is a highly effective risk-mitigation strategy. The data reveals that applicants who would otherwise be rejected due to low individual income achieve viable approval rates when their total household income (TotalIncome) is evaluated via a co-applicant.

